TECHNISCHE UNIVERSITÄT MÜNCHEN
Fakultät für Elektrotechnik und Informationstechnik
Lehrstuhl für Datenverarbeitung
PD Dr. Martin Kleinsteuber

Information Retrieval in High Dimensional Data

Lab #6, 23.11.2017

Principal Component Analysis

Task 1. In this task, we will once again work with the MNIST training set as provided on Moodle. Choose three digit classes, e.g. 1, 2 and 3 and load `N=1000` images from each of the classes to the workspace. Store the data in a normalized matrix `X` of shapetype `double` and size `(784,3*N)`. Furthermore, generate a color label matrix `C` of dimensions `(3*N,3)`. Each row of `C` assigns an RGB color vector to the respective column of X as an indicator of the digit class. Choose `[0, 0, 1]`, `[0,1, 0]` and `[1, 0, 0]` for the three digit classes.

a) Compute the row-wise mean `mu` of `X` and subtract it from each column of `X`. Save the results as `X_c`.

b) Use `np.linalg.svd` with `full_matrices=False` to compute the singular value decomposition `[U,Sigma,VT]` of `X_c`. Make sure the matrices are sorted in descending order with respect to the singular values.

c) Use `reshape` in order to convert `mu` and the first three columns of `U` to `(28,28)`-matrices. Plot the resulting images. What do you see?

d) Compute the matrix `S=np.dot(np.diag(Sigma),VT)`. Note that this yields the same result as `S=np.dot(U.T,X_c)`. The `S` matrix contains the `3*N` scores for the principal components `1` to `784`. Create a 2D scatter plot with `C` as its color parameter in order to plot the scores for the first *two* principal components of the data.

Task 2. In this task, we consider the problem of choosing the number of principal vectors. Assuming that $\mathbf{X} \in \mathbb{R}^{p \times N}$ is the centered data matrix and $\mathbf{P} = \mathbf{U}_k \mathbf{U}_k^\top$ is the projector onto the k-dimensional principal subspace, the dimension k is chosen such that the fraction of overall energy contained in the projection error does not exceed ϵ, i.e.

$$\frac{\|\mathbf{X} - \mathbf{P}\mathbf{X}\|_F^2}{\|\mathbf{X}\|_F^2} = \frac{\sum_{i=1}^{M} \|\mathbf{x}_i - \mathbf{P}\mathbf{x}_i\|^2}{\sum_{i=1}^{N} \|\mathbf{x}_i\|^2} \leq \epsilon,$$

where ϵ is usually chosen to be between 0.01 and 0.2.

The MIT VisTex database as provided on Moodle consists of a set of 167 RGB texture images of sizes `(512,512,3)`. Download the ZIP file, unpack it and make yourself familiar with the directory strucutre.

a) After preprocessing the entire image set (converting to normalized grayscale matrices), divide the images into non overlapping tiles of sizes `(64,64)` and create a centered data matrix `X_c` of size `(p,N)` from them, where `p=64*64` and `N=167*(512/64)*(512/64)`.

b) Compute the SVD of `X_c` and make sure the singular values are sorted in descending order.

c) Plot the fraction of energy contained in the projection error[1] for the principal subspace dimensions `0` to `p`. How many principal vectors do you need to retain 80%, 90%, 95% or 99% of the original data energy?

d) Discuss: Can you imagine a scenario, where energy is a bad measure of useful information?

Helpful Python/Numpy functions

`import scipy.misc`	contains `imread`
`import matplotlib.pyplot`	contains plotting functionalities
`glob.glob(str)`	lists all files that contain `str` in the relative path

[1] Note that you do not need to evaluate any norms or projections. All you need is the result of subtask b)